

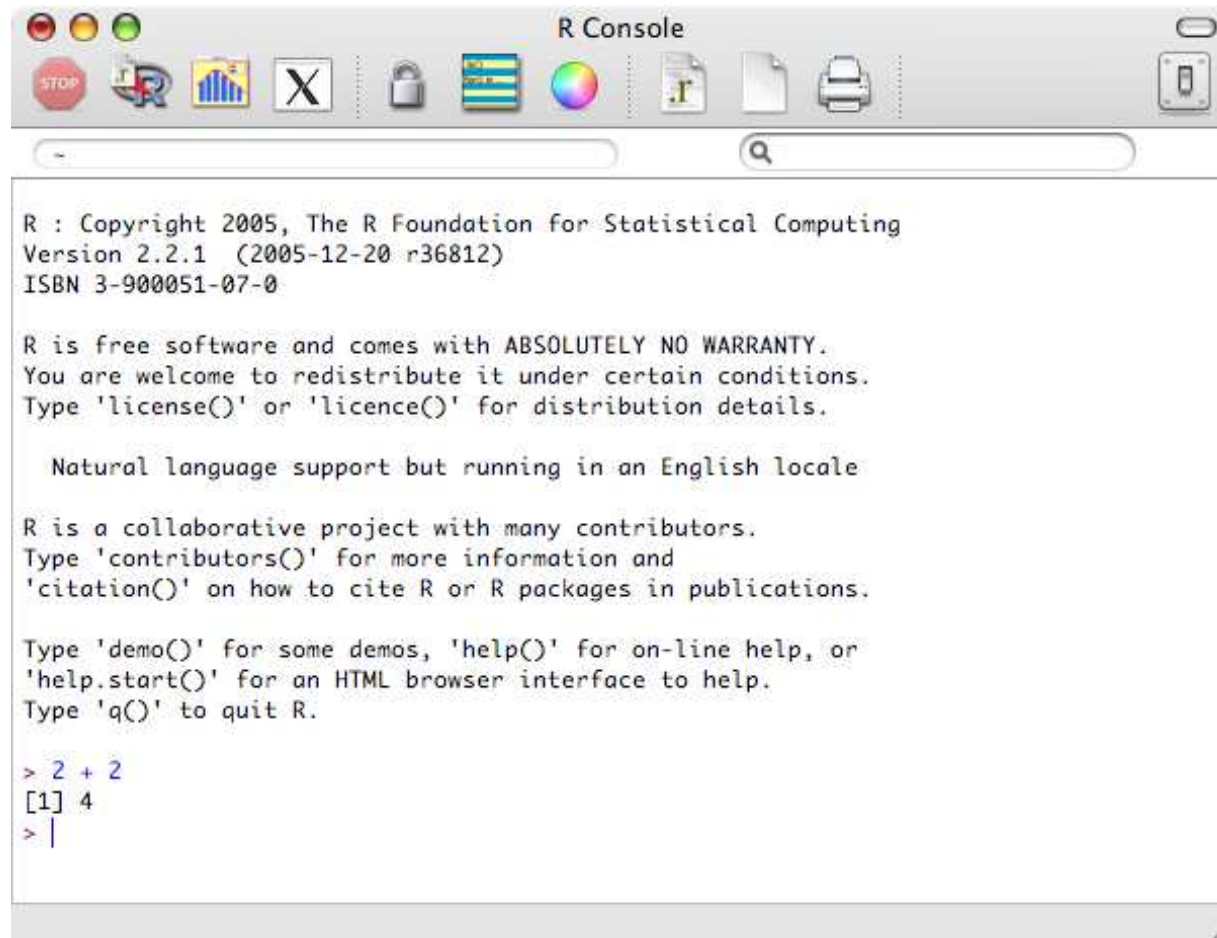


Statistical Computing in R

R is a programming language designed to support data analysis and model building.

- All traditional programming constructs such as expressions, assignments, conditionals, loops, and functions are present.
- A straight-forward object system that supports high-level constructs such as statistical models with all their parameters etc. very nicely.
- Vector arithmetic (very powerful and the preferred way of accomplishing things in R).
- Graphics engine supporting graphical techniques (automatic scatter plots, histograms, etc.)
- Many, many extension modules implementing everything from basic statistics to micro array analysis...and in particular support vector machines.

Interactive R Session



The screenshot shows the R Console window with a standard macOS-style title bar. The window contains the following text:

```
R : Copyright 2005, The R Foundation for Statistical Computing
Version 2.2.1 (2005-12-20 r36812)
ISBN 3-900051-07-0

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

> 2 + 2
[1] 4
> |
```



R Programming

```
> x <- 2
> 2 * x
[1] 4
```

```
> v <- c(1,2,3)
> v
[1] 1 2 3
```

```
> v + 1
[1] 2 3 4
```

```
> w <- v + 1
> q <- w + v
> q
[1] 3 5 7
```

```
> add1 <- function(x) { x + 1; }
> add1
function(x) { x + 1; }
> add1(1)
[1] 2
```

R Programming

```
> addv1
function(v)
{
  y <- c()
  for (x in v) {
    x1 <- x + 1
    y <- c(y, x1)
  }
  y
}

> w
[1] 2 3 4
> addv1(w)
[1] 3 4 5
```

This function performs the same operation as the vector operation $w + 1$. From a performance point of view it is always desirable to use the vector operations, explicit iteration over vector elements is SLOW!



R Data

R has many different ways to represent data:

- vectors
- lists
- arrays/matrices

The most important one (for our purposes) is the *data frame*. A data frame is a two-dimensional data matrix with additional structure.

```
> df <- data.frame(v, w)
> df
  v w
1 1 2
2 2 3
3 3 4
> df$v
[1] 1 2 3
> df$w
[1] 2 3 4
```



Loading Data Frames

We can read comma-separated-value (CSV) files directly into an R data frame.

Here is our mammal training data set represented as a CSV file:

```
Legs, Wings, Fur, Feathers, Mammal
4, no, yes, no, true
2, yes, no, yes, false
4, no, no, no, false
4, yes, yes, no, true
3, no, no, no, false
```

Assume that we saved this into a file called “mammals.csv”, in a directory called “datasets”.

Loading Data Frames

```
> setwd("datasets")
> mammals.df <- read.csv("mammals.csv")
> mammals.df
```

	Legs	Wings	Fur	Feathers	Mammal
1	4	no	yes	no	true
2	2	yes	no	yes	false
3	4	no	no	no	false
4	4	yes	yes	no	true
5	3	no	no	no	false

```
> summary(mammals.df)
```

	Legs	Wings	Fur	Feathers	Mammal
Min.	:2.0	no :3	no :3	no :4	false:3
1st Qu.:	3.0	yes:2	yes:2	yes:1	true :2
Median	:4.0				
Mean	:3.4				
3rd Qu.:	4.0				
Max.	:4.0				

R Built-in Data Frames

For convenience sake, R comes with a number of predefined data frames. One such predefined data frame is the *iris data set*.

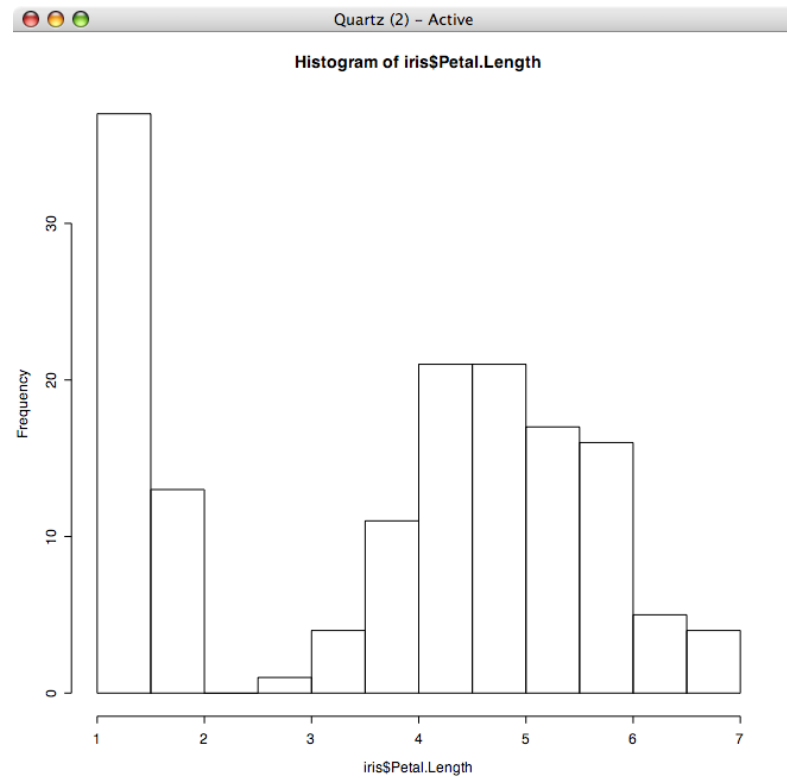

```
> data(iris)
> summary(iris)
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
Min. :4.300	Min. :2.000	Min. :1.000	Min. :0.100	setosa :50
1st Qu.:5.100	1st Qu.:2.800	1st Qu.:1.600	1st Qu.:0.300	versicolor:50
Median :5.800	Median :3.000	Median :4.350	Median :1.300	virginica :50
Mean :5.843	Mean :3.057	Mean :3.758	Mean :1.199	
3rd Qu.:6.400	3rd Qu.:3.300	3rd Qu.:5.100	3rd Qu.:1.800	
Max. :7.900	Max. :4.400	Max. :6.900	Max. :2.500	

We might wish to inspect the data distributions visually as well:

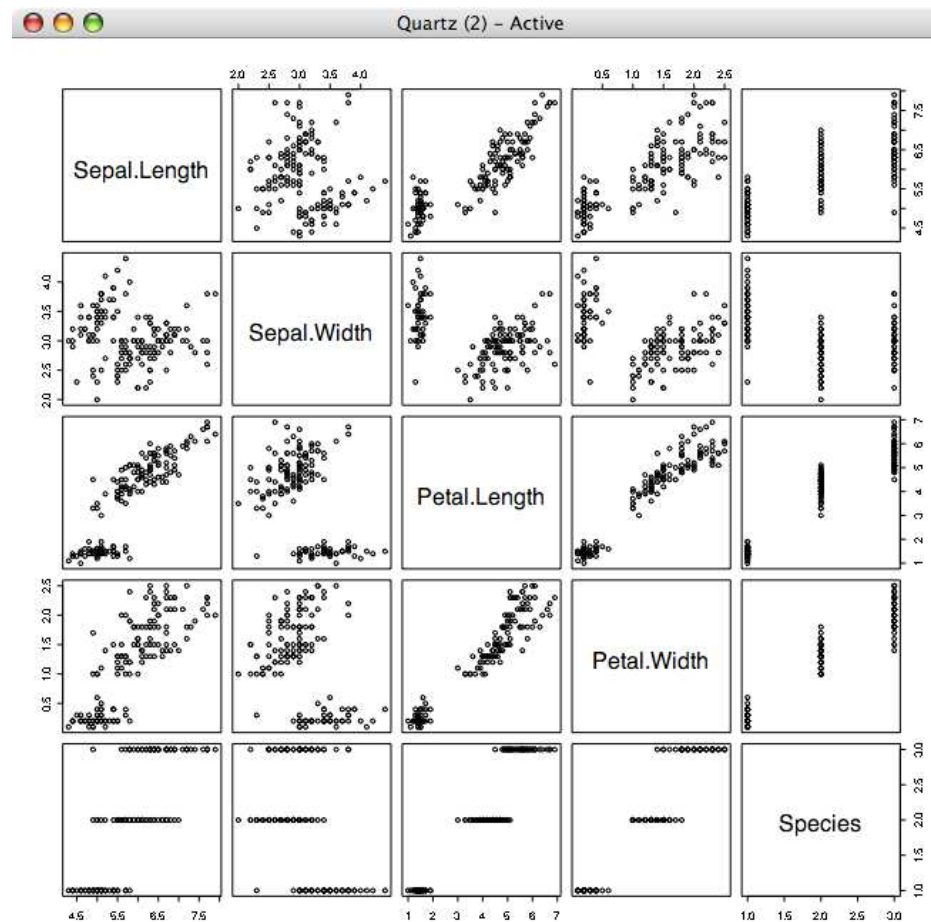
```
> hist(iris$Sepal.Length)
> hist(iris$Petal.Length)
```


R Built-in Data Frames




R Built-in Data Frames

```
> plot(iris)
```



Simple Model Building

```
> attach(iris)
> model <- lm(Petal.Length~Petal.Width)
> plot(Petal.Width,Petal.Length)
> abline(model)
```





Homework

Read Chapter 1 and Appendix B

Do Assignment 1, see website.