



Ensemble Methods

- In statistics and machine learning, ensemble methods use multiple learning algorithms to obtain better predictive performance than could be obtained from any of the constituent learning algorithms.
- Ensemble methods construct a set of classifiers and then classify new data points by taking a weighted *vote* of their predictions.



Ensemble Methods

- Supervised learning algorithms are commonly described as performing the task of searching through a hypothesis space to find a suitable hypothesis that will make good predictions with a particular problem.
- Even if the hypothesis space contains hypotheses that are very well-suited for a particular problem, it may be very difficult to find (a good) one,
 - local minima in ANN's
 - too expensive to fit a decision surface in a highly non-linear situation (SVM) – fitting an SVM with large values of C



Ensemble Methods

- Ensembles combine multiple hypotheses to form a (hopefully) better hypothesis.
- The term *ensemble* is usually reserved for methods that generate multiple hypotheses using the same base learner, *i.e.* decision tree

Note: The broader term of *multiple classifier systems* also covers hybridization of hypotheses that are not induced by the same base learner.



Ensemble Methods

- Evaluating the prediction of an ensemble typically requires more computation than evaluating the prediction of a single model
- Ensembles may be thought of as a way to compensate for poor learning algorithms by performing a lot of extra computation.



Ensemble Methods

- An ensemble is itself a supervised learning algorithm
 - it is trained on labeled data and then used to make predictions on unseen data points
- A trained ensemble represents a single hypothesis,
 - This hypothesis, however, is not necessarily contained within the hypothesis space of the models from which it is built.
 - Thus, ensembles can be shown to have more flexibility in the functions they can represent – we will see that with Random Forests

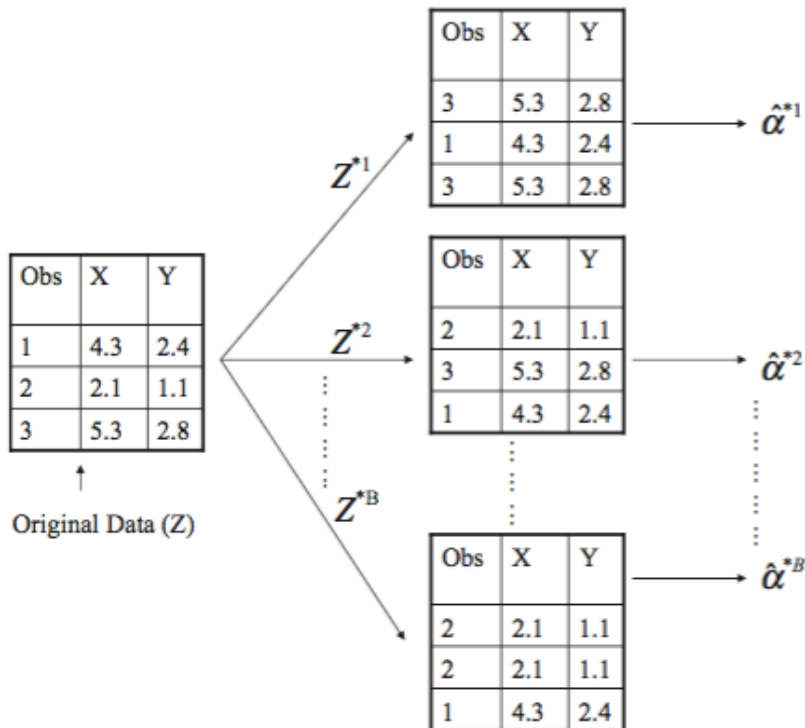


Ensemble Techniques

- Bagged Trees
 - Bootstrap AGGREGatED decision Trees
- Random Forests



Bagged Trees



- Recall the bootstrap
 - resample with replacement – B bootstrap samples - Z^{*k}
- Build B decision trees on the bootstrap samples - $\hat{\alpha}^{*k}$
- Given a point have each tree vote on the prediction,
 - The prediction with the most votes is the prediction of the ensemble.




Bagged Trees

- Notice that due to the resampling with replacement each bootstrap sample represents the input domain slightly differently.
- This means that bootstrapping might eliminate some irregularities from the original data that a single tree might have a difficult time learning/representing.




Bagged Trees



Training data



Bagged Trees




Decision boundary produced
by one tree




Bagged Trees

Banana Set



Decision boundary produced by a second tree


Banana Set




Decision boundary produced by a third tree



Bagged Trees



Three trees and final boundary overlaid



Final result from bagging all trees.



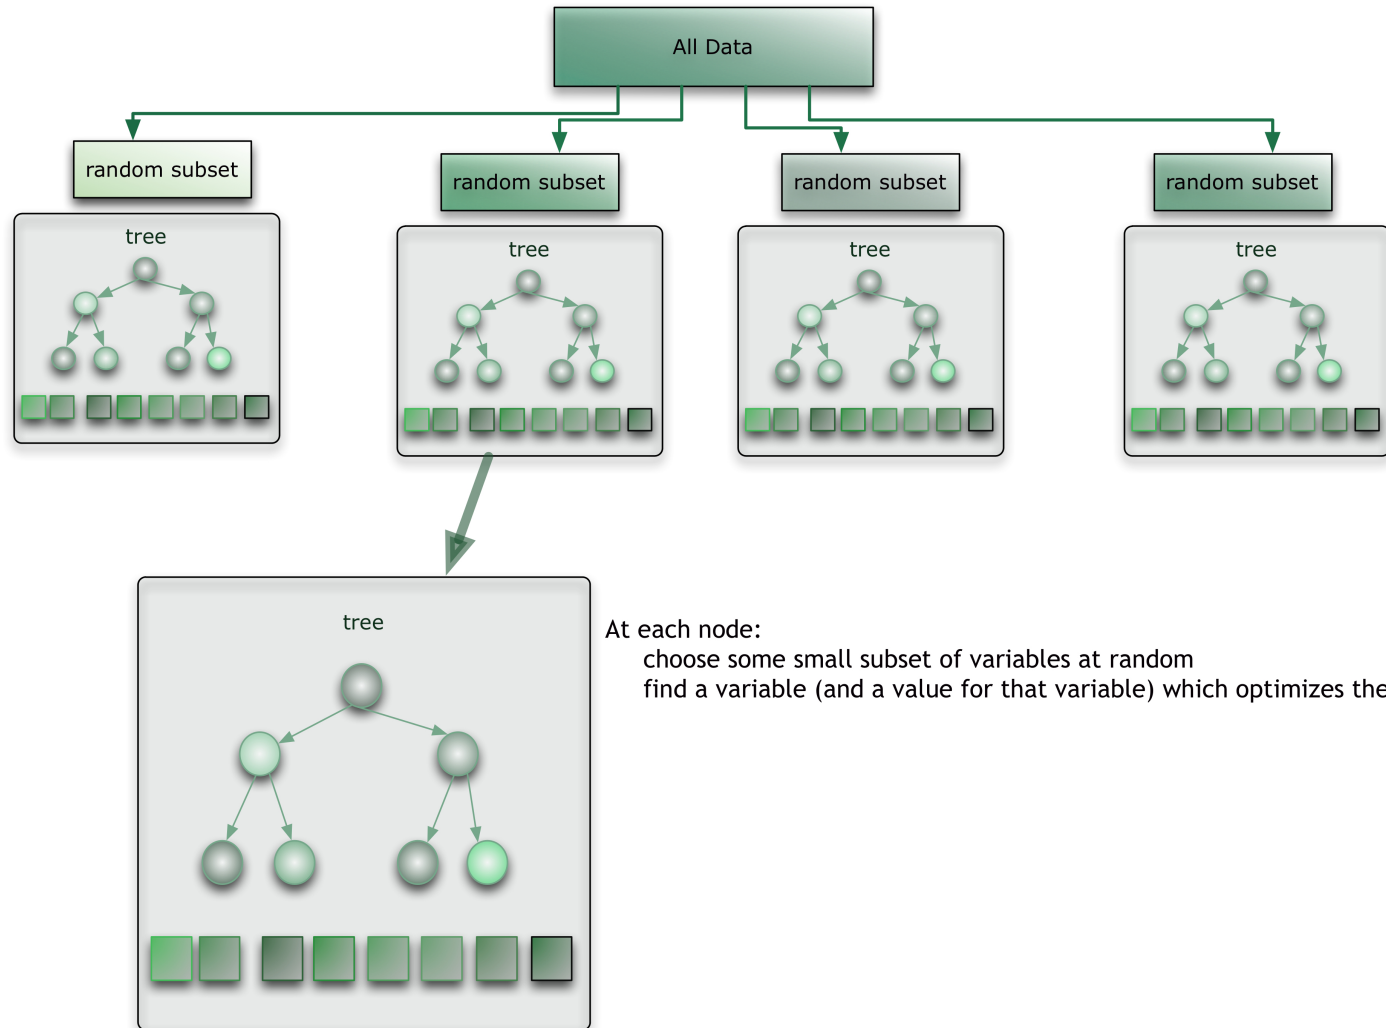
Random Forests

- Very similar to Bagged Trees
- One big difference: the attribute evaluated at a split is drawn from a *random* subset of all possible attributes



source: <https://citizennet.com/blog/2012/11/10/random-forests-ensembles-and-performance-metrics/>

Random Forests





Random Forests

Each tree is grown as follows:

- Create a bootstrap sample from the original data. This sample will be the training set for growing the tree.
- For M input variables/attribute, choose a number $m \ll M$,
 - *at each node, m variables are selected at random out of the M and the best split on these m is used to split the node.* The value of m is held constant during the forest growing.
- Each tree is grown to the largest extent possible. There is no pruning.



Random Forests

ID3(\mathbf{D}, \mathbf{X}) =

Let T be a new tree

If all instances in \mathbf{D} have same class c

Label(T) = c ; Return T

If $\mathbf{X} = \emptyset$ or no attribute has positive information gain

Label(T) = most common class in \mathbf{D} ; return T

$X \leftarrow$ attribute with highest information gain from m randomly selected attributes from \mathbf{X}

Label(T) = X

For each value x of X

$\mathbf{D}_x \leftarrow$ instances in \mathbf{D} with $X = x$

If \mathbf{D}_x is empty

Let T_x be a new tree

Label(T_x) = most common class in \mathbf{D}

Else

$T_x = \text{ID3}(\mathbf{D}_x, \mathbf{X} - \{X\})$

Add a branch from T to T_x labeled by x

Return T

- At top level generate bootstrap samples \mathbf{D}
- After trees have been grown combine in a majority voting scheme